

WEAK SUPERVISION AND ACTIVE LEARNING WITH TEXT DATA

TRAINING A MULTI-CLASS CLASSIFIER WITH WEAK LABELS
AND IMPROVING IT WITH ACTIVE LEARNING

TEAM 44 -  FSDL 2022

Aleks Hiidenhovi Bernardo García

Kushal Ramaiya Edoardo Abati

Juan Manuel Diego Quintana

Sal El Kafrawy

TABLE OF CONTENTS

01.

INTRODUCTION

02.

DEV ENVIRONMENT SETUP

03.

WEAK SUPERVISION AND
ACTIVE LEARNING

04.

MODEL TRAINING

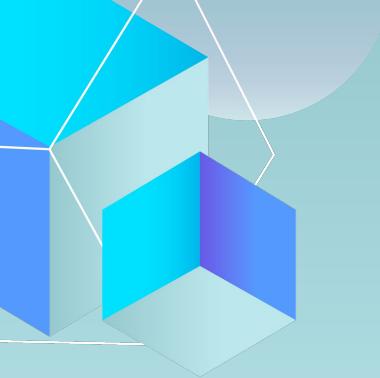
05.

MODEL DEPLOYMENT

06.

FURTHER WORK





01



INTRODUCTION

TRAINING A MULTI-CLASS CLASSIFIER WITH
WEAK LABELS AND IMPROVING IT WITH
ACTIVE LEARNING



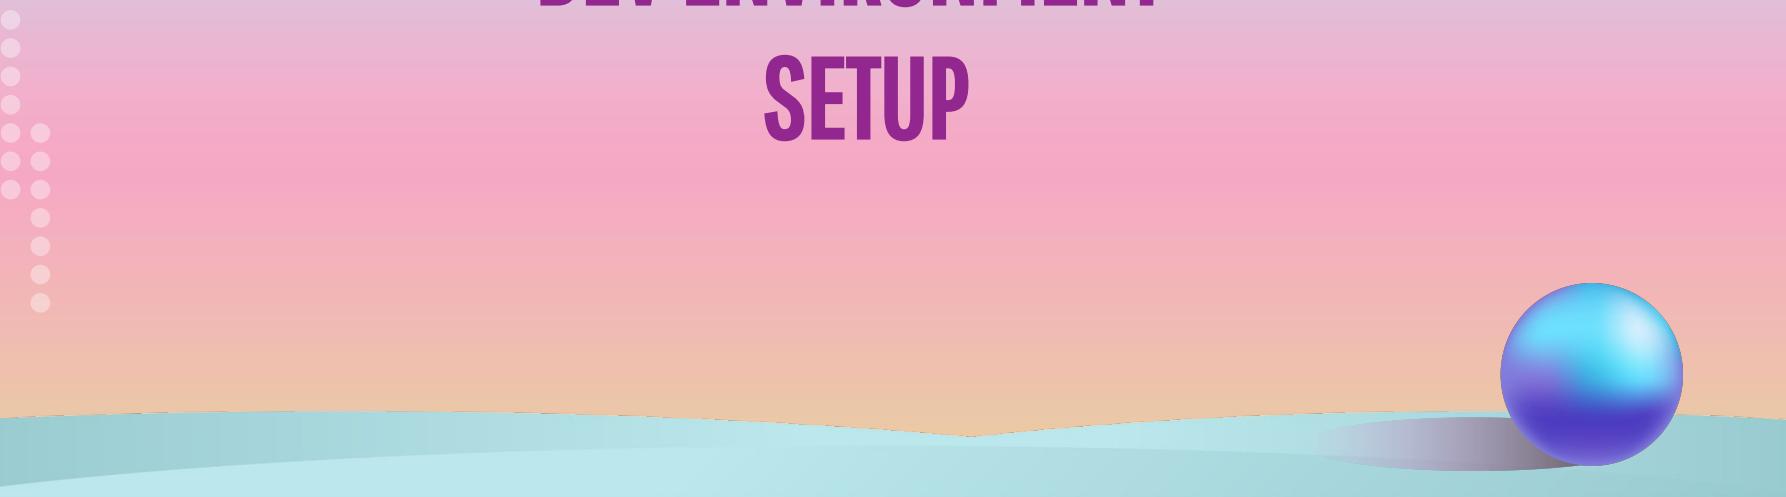
INTRODUCTION

- **News classifier** → Multi-class classification
- We combined **weak supervision** and **active learning** to **minimise** the amount of time for human labelling
- **Data** → We used the AG news dataset without labels for the training stage
- We used an open-source library **Rubrix** to create an efficient workflow
- The classifier consists of a fine-tuned **DistilBERT**:
 - The model was trained with **PyTorch** and **HuggingFace**
 - experiments were tracked with **Weights & Biases**
 - The versions of the model were stored on **W&B's model registry**
 - We deployed the model using **AWS Serverless Lambdas**
 - We implemented a UI using **Streamlit**
- **Code is available at** <https://github.com/EdAbati/fsdl-2022-weak-supervision-project>



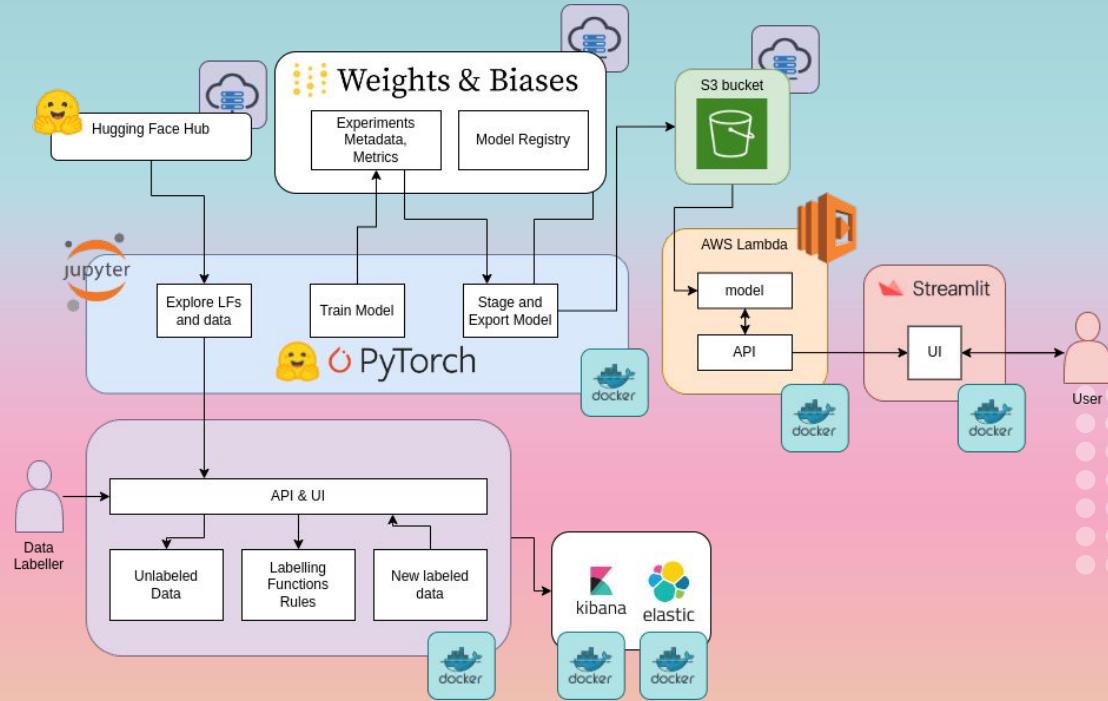
02

DEV ENVIRONMENT SETUP



Development environment

- Six containers on **docker-compose**
- **Streamlit** for UI
- **Development container** with jupyter and dependencies for training and deploying models
- **Rubrix service** for supporting labelling and active learning loop
- **Elasticsearch** and **Kibana** to support Rubrix
- **AWS Lambda API** for serving the model





03

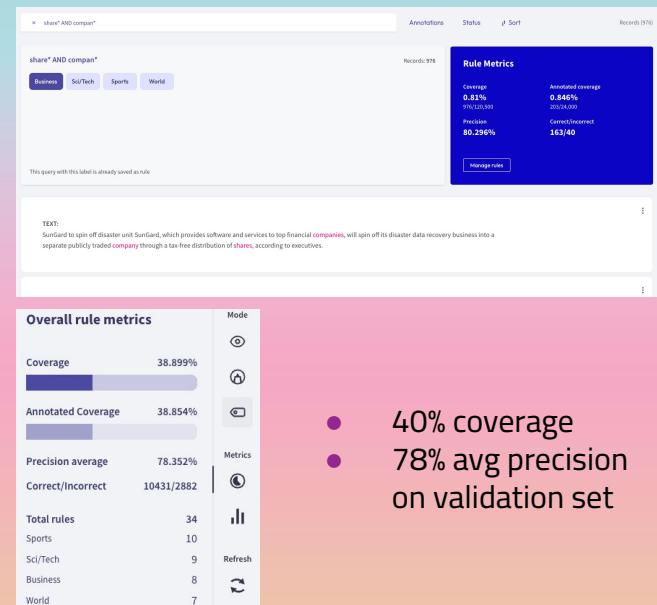
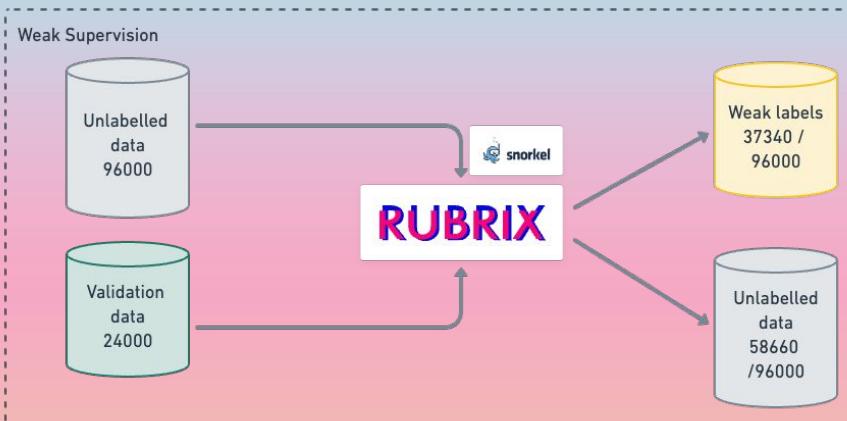


WEAK SUPERVISION AND ACTIVE LEARNING

OPTIMISING HUMAN LABELLING WITH WEAK
SUPERVISION AND ACTIVE LEARNING

WEAK SUPERVISION

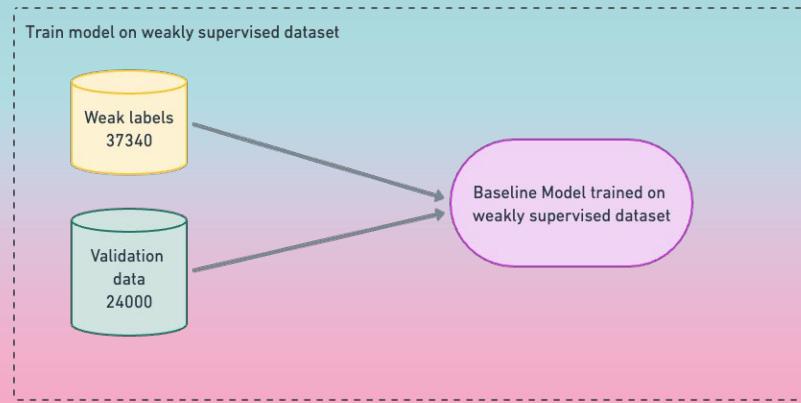
- We wrote **34 labelling functions (LFs)** to create weak labels → Mainly **Pattern Matching LFs**
 - **4 classes:** World, Sports, Sci/Tech and Business
- **Rubrix UI** was very helpful to check the performance of the LFs on the validation set*
- We denoised the weak labels with **Snorkel's label model** to obtain probabilistic labels for the final weakly supervised dataset



*We used 20% of the original AG news dataset as validation set to simulate a real world scenario where this data would have been human-labelled

ACTIVE LEARNING LOOP

- Model trained on **weakly supervised dataset** → **No need to label a training dataset** to initialise the active learning loop



- Active learning** allows us to identify data points that **are confusing the model (uncertainty sampling)** or **are novel to the model (diversity sampling)**, thus **optimising the amount of data that needs to be human-labelled**

Uncertainty Sampling Comparison

Least Confidence

- Picks samples where the **most likely class has the lowest probabilities**
- **Benefits:** very simple and intuitive
- **Shortcomings:** this method is **not very robust**, and doesn't provide the best signal for uncertainty

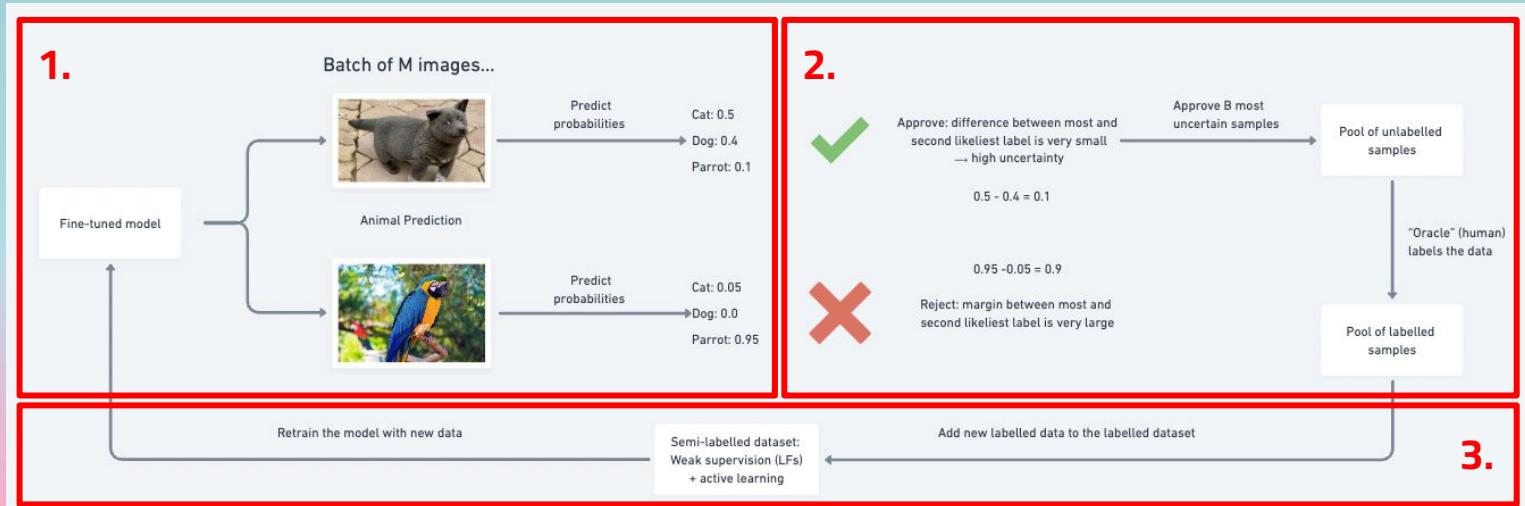
Margin Sampling

- Picks samples where the **two most likely classes have similar probabilities**
- **Benefits:** also **simple** and **more robust** than least confidence
- **Shortcomings:** still only uses the two most probable classes, other classes don't have an effect on uncertainty

Entropy Sampling

- **Averages the uncertainty** across all the classes
- **Benefits: robust**, averages the uncertainty throughout all the classes
- **Shortcomings: math** behind entropy sampling is a bit **more involved** compared to the other two methods

Active Learning Loop with Margin Sampling



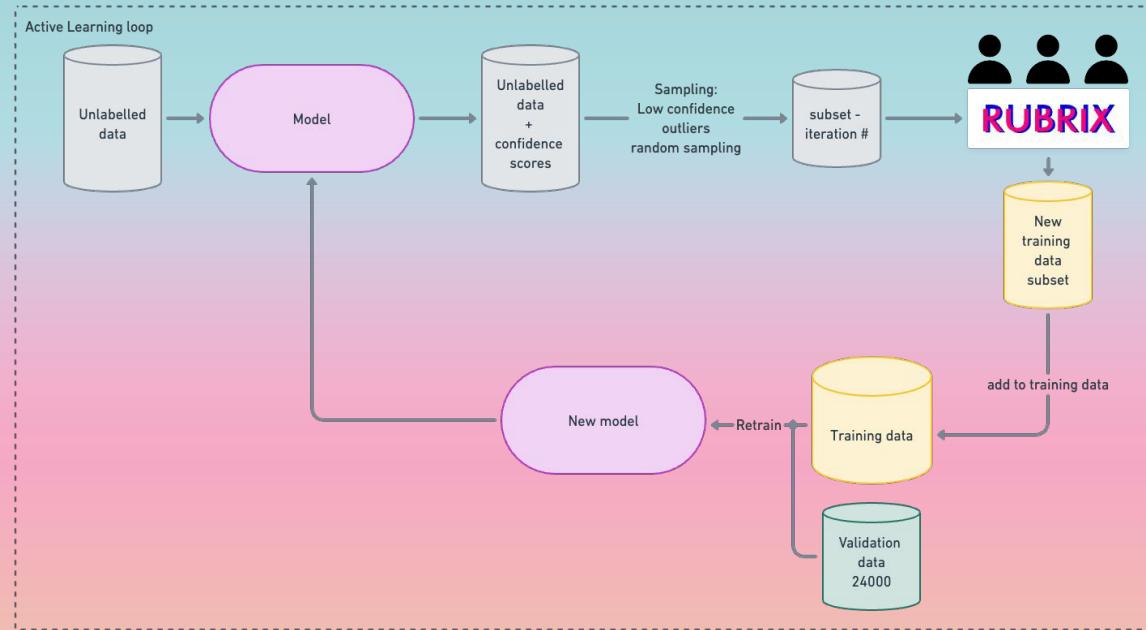
Step 1: Take a pool of M images, use the model to predict the probabilities for the images

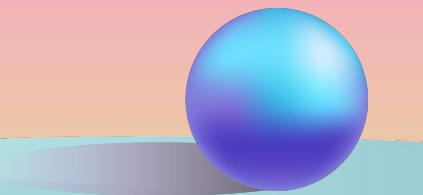
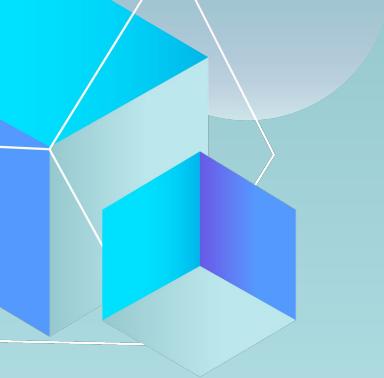
Step 2: Calculate uncertainty scores for all n images, choose B most uncertain samples for human labeling

Step 3: Add newly labelled samples to your training set and retrain your model

ACTIVE LEARNING LOOP - IMPLEMENTATION

1. Model is used to **get confidence scores** from the pool of unlabeled data
2. A new **subset of X samples** is created based on **uncertainty sampling**:
Margin sampling
3. The **new subset** is **human-labeled**
4. The new subset is **added to the training set** and the **model retrained** to obtain a new model
5. The new model is now used to **get confidence scores** and **select samples** from the remaining unlabeled data





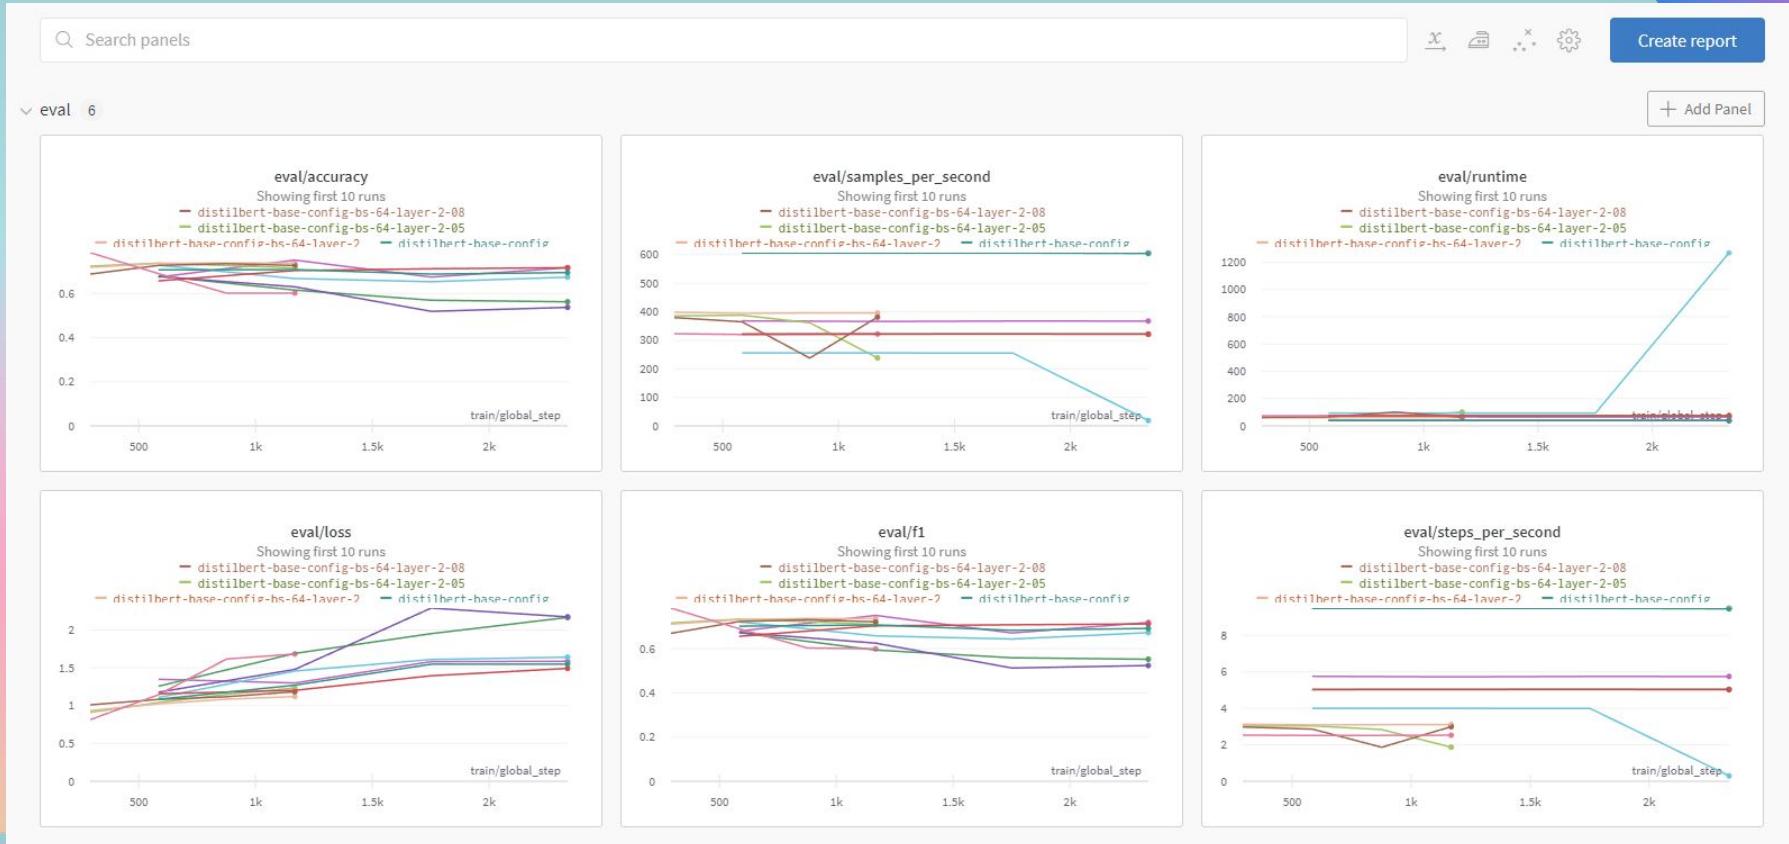
04.

MODEL TRAINING

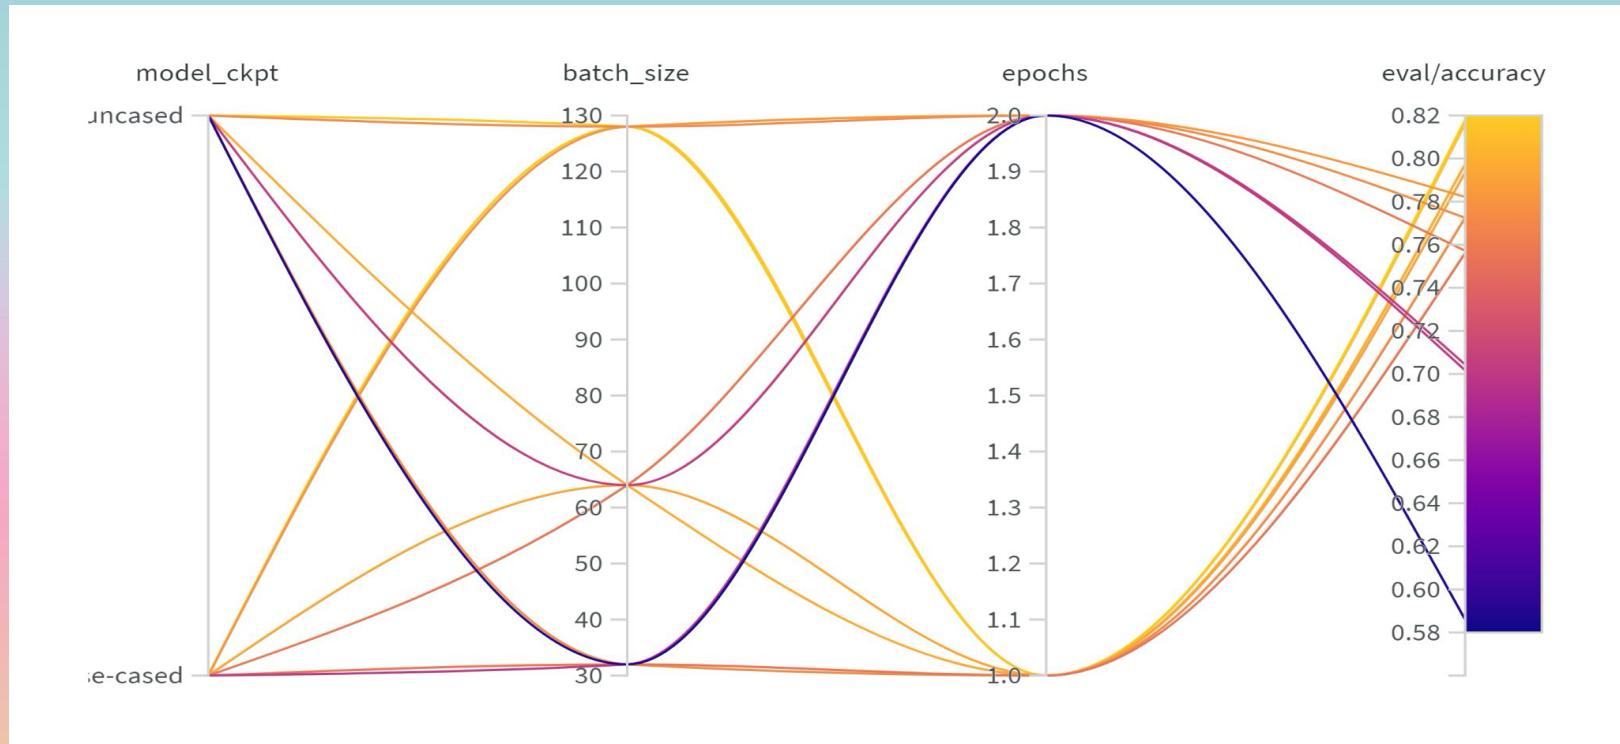
Model Training Loop

- **Natural Language Processing Text MultiClass Classifier**
- Loaded data with weak labels hosted on 😊 Datasets, generated using Rubrix
- The data is **tokenized** using the **subword tokenization** technique.
- We load the model using `.from_pretrained` passing the name of the model as a parameter which loads from 😊 Models.
- Each run is logged into Weights & Biases, where we can compare different runs and their accuracy.
- Further W&B Sweeps are also implemented to automate hyperparameter search along with interactive dashboards.

WandB Individual Runs Combined



Weight And Biases Sweep (Example Run)



https://wandb.ai/team_44/huggingface/sweeps/phsok7we?workspace=user-diegoquintana



05.

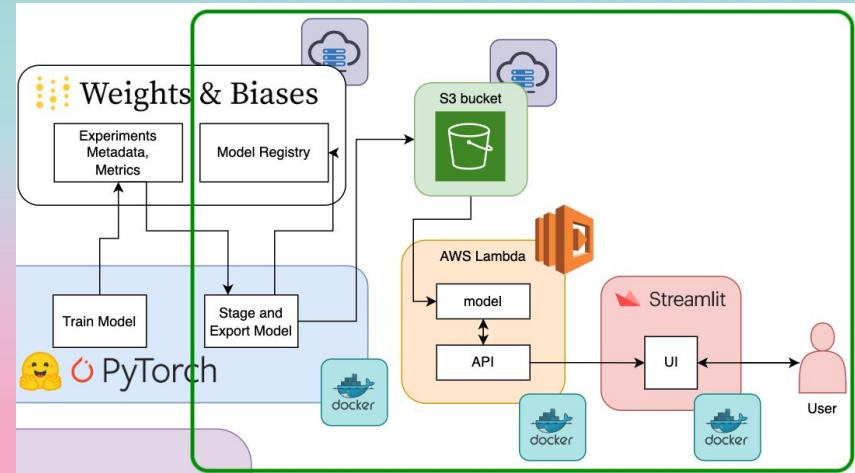
MODEL DEPLOYMENT

HOW THE MODEL WAS DEPLOYED AND
DEPLOYMENT STRATEGY



Model deployment

- The **Weight and Biases model registry** is used to register models and flag the one that is 'production-ready'
- The model selected for deployment is converted into **TorchScript** and uploaded to **S3 Bucket**
- The model inference runs in a **AWS Lambda** function
- User can query the model using a **Streamlit** app. This send requests to the model via the AWS Lambda API endpoint.



User Interface

- [Streamlit link](#)

News Classifier App

Classify a news article!

Write a news article headline in the prompt below.

The app will classify it in one of 4 different categories: World, Sport, Business or Sci/Tech.

Type Here

Wall St. Bears Claw Back Into the Black (Reuters) Reuters - Short-sellers, Wall Street's dwindling band of ultra-cynics, are seeing green again.

Submit

Prediction

Business 📈

Original Text

Wall St. Bears Claw Back Into the Black (Reuters)
Reuters - Short-sellers, Wall Street's
dwindling band of ultra-cynics, are seeing green
again.

Prediction Probability

	label	score	prob
0	Business 📈	0.8468	84.68%
1	Sci/Tech 🚀	0.1326	13.26%
2	Sports 🏆	0.0156	1.56%
3	World 🌎	0.0050	0.50%

score

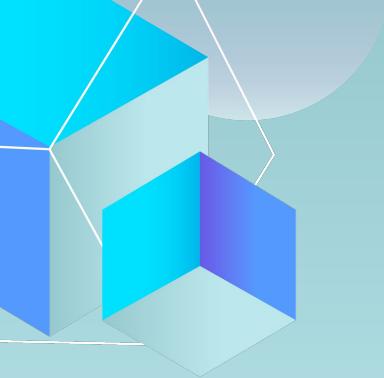
label

label	score
Business 📈	0.8468
Sci/Tech 🚀	0.1326
Sports 🏆	0.0156
World 🌎	0.0050

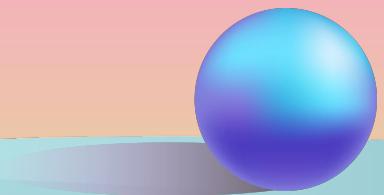
Model deployment

- A model is stored in the **W&B model registry** with the label **prod**
- After a new model is trained, its accuracy ("**eval/f1**") **score is compared against the prod model on the test set***
- **If the score is higher** than the production model, the **challenger model is registered as prod**
- A **traced version** of the model and **its metadata** is uploaded to the S3 bucket
- The **lambda container is rebuilt** so that it fetches the new model

*We used the test set of the original AG news dataset to simulate a real world scenario where this data would have been human-labelled



06



FURTHER WORK

HOW TO IMPROVE IT?

FURTHER WORK

- **Weak supervision** → More comprehensive EDA to build more domain knowledge and create more sophisticated Labeling Functions (LFs)
- **Active learning strategy** → Use Entropy Sampling instead of Margin Sampling or combine different active learning strategies. E.g. 70% margin sampling + 20% Outliers + 10% Random sampling
- **Model training and deployment** → Automate model retraining and deployment after each iteration of the active learning loop
- **Model monitoring** → Come up with relevant projections for monitoring the performance of the news classifier and push them to Gantry
- **Product & UI/UX** → Build a simple product where the model could be useful for people

THANKS!

Do you have any questions?



CREDITS: This presentation template was created by [Slidesgo](#),
including icons by [Flaticon](#), and infographics & images by
[Freepik](#)